

A Byte-Level Multilingual News Recommender for Microcontrollers: NAS, Micro-NAS and Binarized Micro-NAS across Flash, RAM and Energy

Abstract. On-device news recommendation promises privacy, since no behavioural data leaves the device, but it faces a cold-start problem and a hard memory wall: classic neural news recommenders carry a word-embedding table of tens to hundreds of megabytes, far beyond a microcontroller’s flash. We present a tiny, **language-agnostic** news encoder that distils a frozen multilingual sentence-transformer teacher into a **byte-level char-CNN student** whose entire vocabulary is the 256 UTF-8 byte values, eliminating per-language tables and serving all 14 languages of xMIND with one model under 2 MB. We study three architecture-search regimes (**NAS**, **Micro-NAS** (MCU-constrained), and **binarized Micro-NAS**) across three precisions (**FP32** / **INT8** / **Binary**), reporting a full matrix of ranking quality (AUC, MRR, nDCG@5/10 on MIND) against **flash footprint, RAM, and energy** targeting an RTX 5070 $\rightarrow$ Raspberry Pi 5 $\rightarrow$ STM32H7 path (laptop-measured; on-device boards are future work). INT8 Micro-NAS reaches **0.610** AUC at **204 KB**, matching a full NRMS word-embedding baseline (0.607) at a fraction of the size, while unconstrained INT8 NAS reaches **0.647** AUC at 790 KB (+0.04 over NRMS, with no vocabulary table and full multilingual coverage); binarization cuts the footprint to **186–239 KB** at 0.52–0.55 AUC. We release the full pipeline and exportable cold-start prior and content encoder for edge deployment.

Keywords: Edge AI · Neural Architecture Search · Binarized Neural Networks · Multilingual News Recommendation · TinyML · Knowledge Distillation

1 Introduction

Privacy-first reading apps increasingly personalize *on-device*, where the user model never leaves the phone or edge node. Two obstacles dominate. (i) **Cold start**: a per-user model initialized from zero gives poor early recommendations. (ii) **The memory wall**: state-of-the-art news recommenders such as NRMS [23] and NAML encode article titles with a word-embedding matrix (e.g., GloVe-300d [13] over a 30k–100k vocabulary = 37–120 MB in FP32), which cannot fit a microcontroller’s flash; multiplying this per language for a multilingual app is hopeless.

We target both. We pretrain, *fully offline*, a population cold-start prior and a compact content encoder on MIND [24] and its multilingual extension xMIND [5], then compress to the edge. The encoder is a **byte-level char-CNN** distilled [17] from a frozen multilingual teacher: because its only symbols are the 256 UTF-8 bytes, a *single* model handles all 14 xMIND languages with no per-language table, which removes the memory wall and makes all-language evaluation inexpensive.

Our contribution is a systematic **architecture** $\times$ **precision** study for this encoder: NAS, Micro-NAS (with hard STM32H7 flash/RAM/MAC constraints), and binarized Micro-NAS, each at FP32, INT8, and Binary, measured on flash, RAM, and energy across three deployment tiers. This extends our group’s prior NAS-for-BNN work [14] from vision (WakeVision) to multilingual text recommendation.

Contributions. (1) A *language-agnostic byte-level* news encoder, distilled from a frozen multilingual teacher [17, 20], that removes the word-embedding memory wall and serves 14 xMIND languages plus English with a single sub-2 MB model. (2) A systematic *architecture* $\times$ *precision* benchmark (NAS / Micro-NAS / binarized- μ NAS at FP32/INT8/Binary) reporting ranking quality against flash, RAM, and energy on a laptop $\rightarrow$ Pi 5 $\rightarrow$ STM32H7 path. (3) A distillation ablation isolating the gain of distilled initialisation (+3.3 AUC at the Micro-NAS scale), with released cold-start prior and content encoder for deployment.

Novelty. To our knowledge this is the first byte-level, vocabulary-free news recommender studied under a *joint* architecture-search and precision sweep for microcontroller-class targets: prior on-device NAS/BNN work targets vision (e.g. [9, 10, 14]), and existing multilingual news recommenders [5] are server-scale with per-language sub-word vocabularies.

2 Related Work

Neural news recommendation. NRMS [23], NAML [21], LSTUR [1], NPA [22], and DKN [18] established attention- and knowledge-based title/user encoders on MIND [24]; the standard task is impression-level click ranking (AUC/MRR/nDCG), with published large-test AUC in Table 1. All rely on large word-embedding (or entity) tables, which is the component we remove.

Tokenization-free / compact encoders. Byte- and character-level models [2] and feature hashing [6] drop the vocabulary table; byte-level models scale competitively and are robust in multilingual / noisy settings (ByT5 [25], MambaByte [19]), and multilingual distillation [17] aligns translations to a shared space. On-device recommendation is an active area [26, 27]; we bring byte-level encoding + teacher distillation to the on-device news-rec setting under a joint NAS $\times$ precision study.

Table 1. Published MIND (large, test) AUC of standard neural baselines (numbers from [24]). Context only: our study is on MINDsmall dev, so these are an external upper-bound reference, not a same-split comparison.

	DKN [18]	NPA [22]	NAML [21]	LSTUR [1]	NRMS [23]
AUC	0.646	0.667	0.669	0.677	0.678

TinyML NAS and BNNs. MCUNet [9] and μ NAS [8] search under MCU memory budgets; NAS-BNN [10] searches binary networks; XNOR-Net [15] established 1-bit training with full-precision first/last layers. Our group recently applied NAS-BNN to WakeVision at this venue [14]; here we bring the same NAS $\times$ precision methodology to on-device multilingual recommendation.

3 Method

3.1 Task

Impression-level click ranking: given a user’s clicked history and candidate articles, score candidates. Metrics: AUC, MRR, nDCG@5, nDCG@10.

3.2 Language-Agnostic Content Encoder (Teacher $\rightarrow$ Student)

A frozen **paraphrase-multilingual-MiniLM-L12-v2** teacher [20, 16] produces a 384-d *anchor* embedding for each article’s English title. Since xMIND titles are parallel translations, every language’s translation is distilled to the **same** anchor [17]. The student is a **byte-level char-CNN**: an embedding over 256 UTF-8 byte values, depthwise-separable 1-D conv blocks [4], masked mean pooling, and a linear head, trained with a Matryoshka [7] cosine loss so a truncated prefix is usable on-device. No per-language vocabulary exists.

3.3 Recommender

News vectors from the student feed an additive-attention user encoder over click history; scores are user $\cdot$ candidate dot products, trained NRMS-style with one positive and K negatives (softmax cross-entropy), optimised with AdamW [11].

3.4 Architecture Search (Three Arms)

One search space over {channels, depth, out_dim}, evaluated by distillation quality (cosine to anchors) under footprint feasibility, via an evolutionary supernet-free search:

- **NAS**: FP32, loose (no MCU constraint); accuracy ceiling.

- **Micro-NAS**: INT8, hard constraints (flash ≤ 256 KB, RAM ≤ 128 KB, MACs ≤ 2 M; deliberately tighter than the STM32H7 ceiling so they bind at this small text-encoder scale).
- **Binarized Micro-NAS**: Binary, same hard constraints.

Precision is the controlled second axis of the matrix.

3.5 Precision (FP32 $\rightarrow$ INT8 $\rightarrow$ Binary)

Simulated quantization with a straight-through estimator: per-output-channel symmetric INT8, and $\text{sign} \times \text{scale}$ binarization. Following XNOR-Net [15], the byte embedding and first/last projections stay full-precision; only inner conv/linear weights are quantized. Both PTQ and short QAT are supported.

3.6 Footprint & Energy

Params and MACs via `thop`; model size = $\text{params} \times \text{bytes} / \text{precision}$; energy proxy = $\text{MACs} \times \text{per-op energy}$ (Horowitz 45 nm [3]), labelled relative. Latency is *measured* on the laptop (CUDA-synced timing on the GPU; `onnxruntime` on the CPU). On-device Pi 5 and STM32H7 latency/energy (`onnxruntime` + INA219; X-CUBE-AI cycles) are **future work**; those boards are not yet available, so no Pi 5/STM32H7 numbers are reported.

4 Experimental Setup

Datasets. MINDsmall [24]: 50k users; train 156,965 / dev 73,152 impressions; 51,282 articles. xMINDsmall [5]: 14 languages (zho, fin, grn, hat, ind, jpn, kat, ron, som, swi, tam, tha, tur, vie), NLLB-3.3B translations [12]. xMIND supplies translated title/abstract keyed by news id; MIND behaviours are reused. Multilingual evaluation is therefore **cross-lingual transfer over identical English impressions**, not native-language behaviour.

Why MINDsmall? We use MINDsmall (50k users) rather than MINDlarge (1M users, 24M impressions, 161k articles) because the contribution is on-device footprint/energy and multilingual coverage, not leaderboard accuracy. The full study (three arms, three precisions, evolutionary NAS, and 15-language evaluation) already takes 8.2 h on a single laptop GPU on MINDsmall and would be roughly an order of magnitude longer on MINDlarge; the latter is left to future work (Sect. 6).

Protocol. We train on MINDsmall train and evaluate on MINDsmall dev (small has no public test labels). *Our reproduced NRMS baseline is trained and evaluated on the same MINDsmall split as our models*, so the accuracy comparison is like-for-like; the published MINDlarge-test NRMS (AUC 0.6776) [24] is cited only as an external reference, not a same-split comparison. Fixed seed; SHA256-pinned data.

Hardware & measurement. All training and architecture search run on a single laptop: Intel Core Ultra 9 275HX (24 cores), 32 GB RAM, NVIDIA GeForce

Table 2. Deployment and measurement platforms.

Platform	Key spec	Latency / energy method
Laptop (train+search)	RTX 5070 Laptop GPU, Core Ultra 9 275HX, 32GB	CUDA-synced timing; NVML power
Raspberry Pi 5	BCM2712, Cortex-A76 @2.4 GHz	×4 onnxruntime; INA219 (<i>planned</i>)
STM32H743	Cortex-M7 @480 MHz, 2/1 MB flash/SRAM	X-CUBE-AI cycles (<i>planned</i>)

RTX 5070 Laptop GPU (Blackwell, compute capability 12.0, 8 GB), Windows 11, Python 3.13, PyTorch 2.12.1 (CUDA 13.0). The edge targets are a Raspberry Pi 5 (Broadcom BCM2712, quad Arm Cortex-A76 @2.4 GHz; onnxruntime aarch64) and an STM32H743 (Arm Cortex-M7 @480 MHz, 2 MB flash, 1 MB SRAM; ST X-CUBE-AI), summarised in Table 2. GPU power is sampled via NVML; MAC counts via `thop`; the architecture-agnostic energy proxy uses Horowitz 45 nm figures [3]; physical Pi/MCU power uses an INA219 shunt. Binary MCU numbers are analytical (Sect. 6). The full pipeline is driven from a single notebook.

Reproducibility. All runs use a fixed random seed; the datasets are SHA256-pinned (manifest). The code, the driving notebook, the exported artifacts (ONNX encoder + cold-start prior), and the exact dependency lock are released openly.¹ Profiling uses `thop` (MACs) and onnxruntime (CPU latency); MCU footprint follows ST X-CUBE-AI conventions.

5 Results

5.1 Architecture × Precision Matrix (MINDsmall dev)

Searched architectures (channels-depth-out_dim): NAS = 256-4-384 (unconstrained); Micro-NAS = 64-5-384; Bin. μ NAS = 96-2-384. Footprint, MACs, and energy are of the content encoder only.

How to read Table 3. The **arm** is the *search procedure* (each produces a different architecture); the **precision** is the *deployment quantization*. The two are orthogonal, so we run every searched architecture at all three precisions. Each arm’s *native* precision (†: NAS/FP32, Micro-NAS/INT8, Binarized- μ NAS/Binary) is its intended deployment; the off-diagonal cells (e.g. Micro-NAS at Binary) are ablations that isolate the architecture choice from the precision choice.

INT8 is effectively lossless vs. FP32 (often marginally better, a known QAT regularization effect) while cutting energy 7–36× and size ≈ 1.4 –2×. Binarization costs 5–9 AUC points; because the byte embedding and first/last layers stay full-precision (XNOR-Net practice), its size gain over INT8 is modest at this scale. Micro-NAS finds a 204 KB INT8 encoder that fits an aggressive on-device budget at near-baseline accuracy.

¹ Repository URL omitted for double-blind review.

Table 3. Architecture $\times$ precision matrix (MINDsmall dev). Energy proxy via Horowitz 45 nm [3]. **Bold:** best per group; †: arm’s native precision (intended deployment).

Arm	Prec.	AUC	MRR	nDCG@10	Size (KB)	RAM (KB)	Energy (μ J)
NAS	FP32 [†]	0.633	0.326	0.375	1566.8	128	168.7
NAS	INT8	0.647	0.338	0.387	789.8	32	4.65
NAS	Binary	0.588	0.297	0.342	563.1	32	4.25
Micro-NAS	FP32	0.605	0.307	0.352	266.8	32	15.90
Micro-NAS	INT8 [†]	0.610	0.314	0.360	203.9	8	2.30
Micro-NAS	Binary	0.521	0.251	0.294	185.6	8	2.25
Bin. μ NAS	FP32	0.593	0.290	0.336	311.4	48	15.43
Bin. μ NAS	INT8	0.601	0.295	0.344	255.7	12	3.28
Bin. μ NAS	Binary [†]	0.546	0.273	0.317	239.5	12	3.23

Table 4. Measured single-inference content-encoder latency on the laptop (Table 2); batch 1, mean of 200–500 runs. Pi 5 / STM32H7 on-device timings: TODO.

Architecture	GPU (ms)	CPU / ONNX (ms)
NAS (256-4-384)	0.57	0.34
Micro-NAS (64-5-384)	0.65	0.19
Bin. μ NAS (96-2-384)	0.43	0.10

5.2 Latency (measured on the laptop)

Table 4 reports single-inference encoder latency measured on the laptop (GPU via CUDA-synced timing; CPU via onnxruntime, the engine also used on the Pi 5). All configurations are far under the 50 ms on-device budget; CPU latency tracks compute while GPU latency is launch-overhead-dominated at this scale. **Pi 5 and STM32H7 on-device latency/energy are TODO** (boards pending); only laptop numbers are reported.

5.3 Multilingual Transfer (All 14 Languages)

All 14 languages run on one sub-2 MB model with no per-language table (Table 5). Transfer is strongest for Latin-script / higher-resource targets and weakest for low-resource or distinct-script languages (Georgian kat, Finnish fin, Tamil tam), tracking NLLB translation quality.

5.4 Baseline Comparison

Our reproduced NRMS (word-embedding title encoder, 256-d, MINDsmall vocabulary of 25,355 words) has 7.08 M parameters, 91.6% of them in the embedding table, i.e. **~27 MB FP32 / 6.8 MB INT8**. It reaches MINDsmall-dev

Table 5. Cross-lingual AUC of Micro-NAS / INT8 across all 14 xMIND languages plus English (en). Single byte-level model; no per-language vocabulary table.

en	hat	jpn	grn	ron	tur	vie	swl	zho	ind	som	tha	tam	fin	kat
.604	.556	.551	.533	.532	.530	.523	.521	.521	.537	.514	.504	.504	.495	.491

Table 6. Distillation ablation at the Micro-NAS architecture (64-5-384, MINDsmall dev). Distilled initialisation adds +3.3 AUC points at no extra inference cost.

Encoder init	AUC	MRR	nDCG@10
scratch (no distillation)	0.600	0.293	0.343
distilled-init	0.633	0.345	0.390

AUC **0.607** (MRR 0.322, nDCG@10 0.364); the published MINDlarge-test ceiling is AUC 0.6776 [24].

Fairness. (i) *Accuracy is like-for-like*: NRMS and our models use the identical MINDsmall split, impressions, and metric. The byte student **matches NRMS at 204 KB** (INT8 Micro-NAS, 0.610) and **exceeds it at 790 KB** (INT8 NAS, 0.647). (ii) *Footprint*: NRMS is not an edge model, so even at INT8 it is **6.8 MB**, which is 33 $\times$ our 204 KB encoder and still beyond the STM32H7 budget. This gap comes from the per-language word-embedding table; the size axis in Fig. 1 is logarithmic.

Ablation: does distillation help? At the fixed Micro-NAS architecture (64-5-384; MINDsmall dev; 10 epochs; identical seed/data) we compare the news encoder trained from scratch on clicks against one initialised from the distilled student then fine-tuned (Table 6).

Distillation raises the 64-5-384 model to 0.633, close to the unconstrained NAS encoder (0.647) and above the from-scratch model. The improvement comes from the teacher distillation rather than added capacity, which is why we deploy the distilled-init encoder.

5.5 Footprint / Energy Trade-off

Figure 1 shows the Pareto fronts of AUC vs. model size and vs. energy. INT8 points dominate: Micro-NAS / INT8 (0.610 AUC, 204KB, 2.3 μ J) is the on-device sweet spot; NAS / INT8 is the accuracy leader (0.647, 790 KB). Binary points sit lower-left (smallest, lowest energy) but pay an accuracy cost; on Cortex-M7 (no POPCOUNT), their energy advantage over INT8 is small, so they are footprint-motivated rather than speed-motivated. The full study (3 arms $\times$ 3 precisions + NRMS + 15-language eval) ran in 8.2 h on the single laptop of Table 2.




Fig. 1. Pareto fronts: AUC vs. model size (left) and vs. energy proxy (right). Each marker is one arm $\times$ precision configuration. INT8 points (circles) dominate both fronts; the reproduced NRMS baseline is shown as a dashed reference line.




Fig. 2. Per-language AUC of the deployed Micro-NAS/INT8 model across all 14 xMIND languages. A single byte-level encoder with no per-language vocabulary table serves every language.

6 Limitations

- **On-device measurements are pending.** We measure latency on the laptop only (GPU + onnxruntime CPU). Raspberry Pi 5 and STM32H7 boards are not yet available, so their on-device latency/energy are *not measured here* and remain TODO; size/RAM are computed and energy is the Horowitz proxy.
- **Binary-on-MCU is analytical.** No PyTorch $\rightarrow$ 1-bit $\rightarrow$ STM32H7 toolchain exists (X-CUBE-AI’s 1-bit path is Larq/TF only); Cortex-M7 also lacks POPCOUNT, so binary is *footprint-motivated, not latency-motivated*.
- **Cross-lingual transfer only.** xMIND reuses English impressions; results are not native-language click behaviour.
- **MT quality** varies for low-resource languages (grn, hat, som).
- **MINDsmall scale.** MINDlarge experiments are left to future work.

7 Ethics & Data Statement

MIND is used under the Microsoft Research License (non-commercial research); xMIND under CC-BY-NC-SA-4.0. No personal data is collected or generated; all behaviour data is from the public MIND logs. Translations are machine-generated (NLLB-3.3B) and may carry quality artefacts, especially for low-resource languages.

8 Conclusion

A byte-level, language-agnostic encoder removes the embedding-table memory wall and serves 14 xMIND languages plus English (15 evaluated) with a single sub-2 MB model. Across the NAS, Micro-NAS, and binarized Micro-NAS arms at FP32, INT8, and Binary precision, we map the accuracy–footprint–energy trade-off for on-device multilingual news recommendation. INT8 Micro-NAS is the practical sweet spot: 204 KB, 2.3 μ J, AUC 0.610, and all-14-language coverage with no per-language vocabulary table. The pipeline and deployment artefacts (ONNX encoder, cold-start prior) are released openly.

References

1. An, M., Wu, F., Wu, C., Zhang, K., Liu, Z., Xie, X.: Neural news recommendation with long- and short-term user representations. In: Proceedings of ACL (2019)
2. Clark, J.H., Garrette, D., Turc, I., Wieting, J.: CANINE: Pre-training an efficient tokenization-free encoder for language representation. In: Transactions of the ACL (2022), arXiv:2103.06874
3. Horowitz, M.: Computing’s energy problem (and what we can do about it). In: IEEE ISSCC. pp. 10–14 (2014)
4. Howard, A.G., Zhu, M., Chen, B., Kalenichenko, D., Wang, W., Weyand, T., Andreetto, M., Adam, H.: MobileNets: Efficient convolutional neural networks for mobile vision applications. arXiv:1704.04861 (2017)
5. Iana, A., Glavaš, G., Paulheim, H.: MIND your language: A multilingual dataset for cross-lingual news recommendation. In: Proceedings of SIGIR (2024). <https://doi.org/10.1145/3626772.3657867>, xMIND; arXiv:2403.17876
6. Joulin, A., Grave, E., Bojanowski, P., Mikolov, T.: Bag of tricks for efficient text classification. Proceedings of EACL (2017), arXiv:1607.01759
7. Kusupati, A., Bhatt, G., Rege, A., Wallingford, M., Sinha, A., Ramanujan, V., Howard-Snyder, W., Chen, K., Kakade, S., Jain, P., Farhadi, A.: Matryoshka representation learning. In: Proceedings of NeurIPS (2022), arXiv:2205.13147
8. Liberis, E., Dudziak, Ł., Lane, N.D.: μ NAS: Constrained neural architecture search for microcontrollers. Proceedings of EuroMLSys (2021), arXiv:2010.14246
9. Lin, J., Chen, W.M., Lin, Y., Cohn, J., Gan, C., Han, S.: MCUNet: Tiny deep learning on iot devices. In: Proceedings of NeurIPS (2020), arXiv:2007.10319
10. Lin, M., et al.: NAS-BNN: Neural architecture search for binary neural networks. Pattern Recognition (2024), arXiv:2408.15484; code: github.com/VDIGPKU/NAS-BNN

11. Loshchilov, I., Hutter, F.: Decoupled weight decay regularization. In: Proceedings of ICLR (2019), arXiv:1711.05101
12. NLLB Team: No language left behind: Scaling human-centered machine translation (2022), arXiv:2207.04672
13. Pennington, J., Socher, R., Manning, C.D.: GloVe: Global vectors for word representation. In: Proceedings of EMNLP (2014)
14. Pighetti, A., Bellotti, F., Mohammady, S., Lazzaroni, L., Ballout, H., Berta, R.: Exploring neural architecture search (NAS) for binary neural networks (BNNs) in the new WakeVision. In: Applications in Electronics Pervading Industry, Environment and Society (ApplePies), Lecture Notes in Electrical Engineering, vol. 1553 (2026). https://doi.org/10.1007/978-3-032-17174-0_12
15. Rastegari, M., Ordonez, V., Redmon, J., Farhadi, A.: XNOR-Net: Imagenet classification using binary convolutional neural networks. In: Proceedings of ECCV (2016)
16. Reimers, N., Gurevych, I.: Sentence-BERT: Sentence embeddings using siamese BERT-networks. In: Proceedings of EMNLP-IJCNLP (2019)
17. Reimers, N., Gurevych, I.: Making monolingual sentence embeddings multilingual using knowledge distillation. In: Proceedings of EMNLP (2020), arXiv:2004.09813
18. Wang, H., Zhang, F., Xie, X., Guo, M.: DKN: Deep knowledge-aware network for news recommendation. In: Proceedings of WWW (2018)
19. Wang, J., Gangavarapu, T., Yan, J.N., Rush, A.M.: MambaByte: Token-free selective state space model. arXiv:2401.13660 (2024)
20. Wang, W., Wei, F., Dong, L., Bao, H., Yang, N., Zhou, M.: MiniLM: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. In: Proceedings of NeurIPS (2020), arXiv:2002.10957
21. Wu, C., Wu, F., An, M., Huang, J., Huang, Y., Xie, X.: Neural news recommendation with attentive multi-view learning. In: Proceedings of IJCAI (2019)
22. Wu, C., Wu, F., An, M., Huang, J., Huang, Y., Xie, X.: NPA: Neural news recommendation with personalized attention. In: Proceedings of KDD (2019)
23. Wu, C., Wu, F., Ge, S., Qi, T., Huang, Y., Xie, X.: Neural news recommendation with multi-head self-attention. In: Proceedings of EMNLP-IJCNLP, pp. 6389–6394 (2019)
24. Wu, F., Qiao, Y., Chen, J.H., Wu, C., Qi, T., Lian, J., Liu, D., Xie, X., Gao, J., Wu, W., Zhou, M.: MIND: A large-scale dataset for news recommendation. In: Proceedings of ACL, pp. 3597–3606 (2020)
25. Xue, L., Barua, A., Constant, N., Al-Rfou, R., Narang, S., Kale, M., Roberts, A., Raffel, C.: ByT5: Towards a token-free future with pre-trained byte-to-byte models. Transactions of the ACL (2022), arXiv:2105.13626
26. Yin, H., Qu, L., Chen, T., Yuan, W., Zheng, R., Long, J., Xia, X., Shi, Y., Zhang, C.: On-device recommender systems: A comprehensive survey. arXiv:2312.10864 (2024)
27. Yu, Y., Wu, F., Wu, C., Yi, J., Liu, Q.: Tiny-NewsRec: Effective and efficient plm-based news recommendation. Proceedings of EMNLP (2022), arXiv:2112.00944